# ME 759
## High Performance Computing for Engineering Applications
## Assignment 3
## Due Thursday 09/29/2022 at 9:00 PM

Submit responses to all tasks which don't specify a file name to Canvas in a file called assignment3.{txt, docx, pdf, rtf, odt} (choose one of these formats). Submit all plots (if any) on Canvas. Do not zip your Canvas submission.

All *source files* should be submitted in the `HW03` subdirectory on the `main` branch of your GitLab repo. Please use the name `HW03` exactly as shown here (both in terms of capitalization & name). The `HW03` subdirectory should have no subdirectories. For this assignment, your `HW03` folder should contain `task1.cu`, `task2.cu`, `task3.cu`, and `vscale.cu`.

All commands or code must work on *Euler* with only the `nvidia/cuda` module loaded. The commands may behave differently on your computer, so be sure to test on *Euler* before you submit.

Please submit clean code. Consider using a formatter like clang-format.

IMPORTANT: Before you begin, copy any provided files from `Assignments/HW03` directory of the ME759 Resource Repo. Do not change any of the provided files since these files will be overwritten with clean, reference copies when grading.

---

1. Write a C++ program using CUDA in a file called `task1.cu` which computes the factorial of integers from 1 to 8, by launching a GPU kernel with 1 block and 8 threads. Inside the kernel, each thread should use `std::printf` to write out `a!=b` (followed by a newline), where `a` is one of the 8 integers, and `b` is the result of `a!`. (Follow your kernel call with a call to `cudaDeviceSynchronize()` so that the host waits for the kernel to finish printing before returning from `main`.)

   - Compile: `nvcc task1.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std=c++17 -o task1`
   - Run (on Euler, use Slurm sbatch!): `./task1`
   - Expected output (showing only 4 out of the 8 lines expected; lines could be out of order):
     ```
     1!=1
     2!=2
     3!=6
     4!=24
     ```

   - It is ok to publish your sbatch script on Piazza. For this assignment, you will need to ask Slurm to execute your program on a node that has GPU cards

2. Write a C++ program using CUDA in a file called `task2.cu` which does the following:

- From the host, allocates an array of 16 `int`s on the device called `dA`.
- Launches a kernel with 2 blocks, each block having 8 threads.
- Each thread computes `ax+y` and writes the result in one distinct entry of the `dA` array. Here,
  - `x` is the thread's `threadIdx`;
  - `y` is the thread's `blockIdx`;
  - `a` is an integer argument that the kernel takes (so all threads use the same `a`). You need to generate `a` randomly and then call the kernel with it. It is up to you how you generate this random number, one possible approach is described here BestPractice.
- Copies back the data stored in the device array `dA` into a host array called `hA`.
- Prints (from the host) the 16 values stored in the host array separated by a single space each.

How to go about it, and what the expected output looks like:

- Compile: `nvcc task2.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std=c++17 -o task2`
- Run (on Euler, use Slurm sbatch!): `./task2`
- Expected output (followed by newline; yours could be different depending on the random number generation): `0 10 20 30 40 50 60 70 1 11 21 31 41 51 61 71`

3. a) Implement in a file called `vscale.cu`, the `vscale` kernel function as declared and described in `vscale.cuh`. This function should take in two arrays, `a` and `b`, and do an element-wise multiplication of the two arrays: $b_i = a_i \cdot b_i$. In the process, `b` gets overwritten. Each thread should do at most one of the multiplication operations.

   **Example:**

   $$a = [-5.0, 2.0, 1.5], \quad b = [0.8, 0.3, 0.6], \quad n = 3$$

   The resulting `b` array is:

   $$b = [-4.0, 0.6, 0.9].$$

   b) Write a file `task3.cu` which does the following:
   - Creates two arrays of length `n` filled by random numbers[1] where `n` is read from the first command line argument. The range of values for array `a` is [-10.0, 10.0], whereas the range of values for array `b` is [0.0, 1.0].
   - Calls your `vscale` kernel with a 1D execution configuration that uses 512 threads per block.
   - Prints the amount of time taken to execute the kernel in *milliseconds* using CUDA events[2].
   - Prints the first element of the resulting array.
   - Prints the last element of the resulting array.

   How to go about it, and what the expected output looks like:
   - Compile: `nvcc task3.cu vscale.cu -Xcompiler -O3 -Xcompiler -Wall -Xptxas -O3 -std=c++17 -o task3`
   - Run (where `n` is a positive integer): `./task3 n`
   - Example expected output (followed by newline):
     ```
     0.012
     1.3
     2.3
     ```

   c) On an Euler *compute node*, run `task3` for each value $n = 2^{10}, 2^{11}, \cdots, 2^{29}$ and generate a plot `task3.pdf` which plots the time taken by your `vscale` as a function of `n`. Overlay another plot which shows the scaling results when using 16 threads per block.

---

[1]Details about random number generation can be found in random_numbers.md.
[2]Recall the GPU timing section of the document timing.md.